

CSP – 554 Big Data Technologies

Assignment #7

Rutul Mehta

A20476293

- Copy “TestDataGen.class” file to the /home/Hadoop directory
- Run Java TestDataGen command to generate magic number.
- Magic number is 104360.
- This command also generates two separate files name foodratings104360.txt and foodplaces104360.txt files in /home/Hadoop directory.
- Then I copied both the files from /home/Hadoop directory to /user/Hadoop directory.

```

E:::::E M:::::M M:::::M M:::::M R:::::R R:::::R
E:::::EEEEEEEEEE M:::::M M:::::M M:::::M M:::::M R:::::RRRRRRR:::::R
E:::::EEEEEEEEEE M:::::M M:::::M M:::::M R:::::RRRRRRR:::::RR
E:::::EEEEEEEEEE M:::::M M:::::M M:::::M R:::::RRRRRRR:::::R
E:::::E M:::::M M:::::M R:::::R R:::::R
E:::::E EEEEE M:::::M MMM M:::::M R:::::R R:::::R
EE:::::EEEEEEEEEE:::::E M:::::M M:::::M R:::::R R:::::R
E:::::EEEEEEEEEE:::::E M:::::M M:::::M RR:::::R R:::::R
EEEEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRR RRRRRR

[hadoop@ip-172-31-14-124 ~]$ java TestDataGen
Magic Number = 104360
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -ls /
Found 3 items
drwxrwxrwt - hdfs hdfsadmingroup 0 2021-10-15 03:56 /tmp
drwxr-xr-x - hdfs hdfsadmingroup 0 2021-10-15 03:55 /user
drwxr-xr-x - hdfs hdfsadmingroup 0 2021-10-15 03:55 /var
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -ls /user/
Found 5 items
drwxrwxrwx - hadoop hdfsadmingroup 0 2021-10-15 03:55 /user/hadoop
drwxrwxrwx - livy livy 0 2021-10-15 03:55 /user/livy
drwxrwxrwx - root hdfsadmingroup 0 2021-10-15 03:55 /user/root
drwxrwxrwx - spark spark 0 2021-10-15 03:55 /user/spark
drwxrwxrwx - zeppelin hdfsadmingroup 0 2021-10-15 03:55 /user/zeppelin
n
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -ls /user/hadoop/
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -ls /home/hadoop/
ls: '/home/hadoop/': No such file or directory
[hadoop@ip-172-31-14-124 ~]$ pwd
/home/hadoop
[hadoop@ip-172-31-14-124 ~]$ ls
foodplaces104360.txt foodratings104360.txt TestDataGen.class
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -copyFromLocal foodplaces104360.txt /user/
hadoop/
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -copyFromLocal foodratings104360.txt /user/
r/hadoop/
[hadoop@ip-172-31-14-124 ~]$ hadoop fs -ls /user/hadoop/
Found 2 items
-rw-r--r-- 1 hadoop hdfsadmingroup 59 2021-10-15 04:24 /user/hadoop/fo
odplaces104360.txt
-rw-r--r-- 1 hadoop hdfsadmingroup 17457 2021-10-15 04:25 /user/hadoop/fo
odratings104360.txt

```

```
>>> foodratings = spark.read.schema(struct1).csv('hdfs:///user/hadoop/foodratings104360.txt')
>>> foodratings.show()
+-----+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+-----+
|Joy|44|21|9|21|3|
|Sam|13|17|6|8|2|
|Joe|21|7|12|8|5|
|Joe|37|19|10|36|3|
|Jill|20|11|9|9|3|
|Sam|34|2|50|39|3|
|Joe|33|38|18|47|4|
|Sam|45|25|42|33|4|
|Mel|42|30|20|4|2|
|Sam|20|47|8|46|1|
|Jill|34|41|9|46|1|
|Joe|25|1|35|23|2|
|Jill|8|6|20|30|3|
|Mel|25|7|2|16|5|
|Jill|48|9|5|12|5|
|Joy|16|18|3|8|2|
|Joy|8|21|20|46|2|
|Joe|15|2|34|9|2|
|Joe|47|30|5|46|2|
|Joe|7|23|43|25|4|
+-----+-----+-----+-----+-----+
only showing top 20 rows

>>> foodratings.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>>
```

Exercise: 2

- Load the 'foodplaces' file as a 'csv' file into a DataFrame called "foodplaces"

```
>>> struct2 = StructType().add("placeid", IntegerType(), True).add("placename", StringType(), True)
>>> foodplaces = spark.read.schema(struct2).csv('hdfs:///user/hadoop/foodplaces104360.txt')
>>> foodplaces.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> foodplaces.show(5)
+-----+-----+
|placeid|placename|
+-----+-----+
|1|China Bistro|
|2|Atlantic|
|3|Food Town|
|4|Jake's|
|5|Soup Bowl|
+-----+-----+

>>> |
```

Exercise: 3

- Step-A

Register the DataFrames created in exercise 1 and 2 as tables called "foodratingsT" and "foodplacesT"

```
>>> foodratings.registerTempTable("foodratingsT")
>>> foodratings2 = spark.sql("select * from ")
Traceback (most recent call last):
  File "/usr/lib/spark/python/pyspark/context.py", line 278, in signal_handler
    raise KeyboardInterrupt()
KeyboardInterrupt

>>> foodratings2 = spark.sql("select * from foodratingsT")
21/10/15 06:41:09 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.validation is not enabled so recording the schema version 1.2.0
21/10/15 06:41:09 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
21/10/15 06:41:10 WARN ObjectStore: Failed to get database global_temp, returning NoSuchObjectException
>>> foodratingsT.show()
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
NameError: name 'foodratingsT' is not defined
>>> foodratings2.show()
+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+
|Joy|44|21|9|21|3|
|Sam|13|17|6|8|2|
|Joe|21|7|12|8|5|
|Joe|37|19|10|36|3|
|Jill|20|11|9|9|3|
|Sam|34|2|50|39|3|
|Joe|33|38|18|47|4|
|Sam|45|25|42|33|4|
|Mel|42|30|20|4|2|
|Sam|20|47|8|46|1|
|Jill|34|41|9|46|1|
|Joe|25|1|35|23|2|
|Jill|8|6|20|30|3|
|Mel|25|7|2|16|5|
|Jill|48|9|5|12|5|
|Joy|16|18|3|8|2|
|Joy|8|21|20|46|2|
|Joe|15|2|34|9|2|
|Joe|47|30|5|46|2|
|Joe|7|23|43|25|4|
+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```

>>> foodplaces.registerTempTable("foodplacesT")
>>> foodplaces2 = spark.sql("select * from foodplacesT")
>>> foodplaces2.show()
+-----+-----+
|placeid| placename|
+-----+-----+
|      1|  China Bistro|
|      2|    Atlantic|
|      3|   Food Town|
|      4|    Jake's|
|      5|   Soup Bowl|
+-----+-----+
>>> |

```

- **Step-B**

Use a SQL query on the table “foodratingsT” to create a new DataFrame called foodratings_ex3a holding records which meet the following condition: food2 < 25 and food4 > 40. Remember, when defining conditions in your code use maximum parentheses.

```

>>> foodratings_ex3a = spark.sql("select * from foodratingsT where (food2<25) and (food4>40)")
>>> foodratings_ex3a.show()
+-----+-----+-----+-----+-----+-----+
|name|food1|food2|food3|food4|placeid|
+-----+-----+-----+-----+-----+-----+
|Joy|8|21|20|46|2|
|Joy|30|2|30|47|1|
|Joy|22|24|1|46|5|
|Jill|12|20|3|48|3|
|Joe|26|24|25|47|4|
|Joe|17|23|2|50|3|
|Jill|5|9|46|49|4|
|Joe|45|19|18|42|5|
|Jill|26|11|20|47|2|
|Jill|20|21|33|45|2|
|Jill|33|17|44|44|3|
|Joe|21|11|22|41|4|
|Mel|1|5|2|41|1|
|Joe|33|3|4|48|2|
|Mel|27|15|36|48|1|
|Mel|24|10|19|48|4|
|Joe|1|6|30|41|3|
|Mel|7|1|23|43|1|
|Jill|29|21|14|44|5|
|Sam|12|13|11|50|1|
+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

>>> foodratings_ex3a.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)
>>>

```

- **Step-C**

Use a SQL query on the table “foodplacesT” to create a new DataFrame called foodplaces_ex3b holding records which meet the following condition: placeid > 3

```

>>> foodplaces_ex3b = spark.sql("select * from foodplacesT where (placeid>3)")
>>> foodplaces_ex3b.show()
+-----+-----+
|placeid| placename|
+-----+-----+
|      4|    Jake's|
|      5|   Soup Bowl|
+-----+-----+

>>> foodplaces_ex3b.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- placename: string (nullable = true)
>>> |

```

Exercise: 4

- Use a transformation (not a SparkSQL query) on the DataFrame 'foodratings' created in exercise 1 to create a new DataFrame called foodratings_ex4 that includes only those records (rows) where the 'name' field is "Mel" and food3 < 25.

```
>>> foodratings_ex4 = foodratings.filter((foodratings.name=="Mel") & (foodratings.food3<25))
>>> foodratings_ex4.show(5)
+----+-----+
|name|food1|food2|food3|food4|placeid|
+----+-----+
| Mel|   42|   30|   20|    4|     2|
| Mel|   25|    7|    2|   16|     5|
| Mel|    9|   22|    9|   14|     4|
| Mel|   39|    4|    9|   26|     1|
| Mel|   14|    4|   24|   27|     2|
+----+-----+
only showing top 5 rows

>>> foodratings_ex4.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>>
```

Exercise: 5

- Use a transformation (not a SparkSQL query) on the DataFrame 'foodratings' created in exercise 1 to create a new DataFrame called foodratings_ex5 that includes only the columns (fields) 'name' and 'placeid'

```
>>> foodratings_ex5 = foodratings.select(foodratings["name"],foodratings["placeid"])
>>> foodratings_ex5.show(5)
+----+-----+
|name|placeid|
+----+-----+
| Joy|      3|
| Sam|      2|
| Joe|      5|
| Joe|      3|
| Jill|     3|
+----+-----+
only showing top 5 rows

>>> foodratings_ex5.printSchema()
root
 |-- name: string (nullable = true)
 |-- placeid: integer (nullable = true)

>>> |
```

Exercise: 6

- Use a transformation (not a SparkSQL query) to create a new DataFrame called ex6 which is the inner join, on placeid, of the DataFrames 'foodratings' and 'foodplaces' created in exercises 1 and 2

```
>>> ex6 = foodratings.join(foodplaces, on=['placeid'], how='inner')
>>> ex6.printSchema()
root
 |-- placeid: integer (nullable = true)
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placename: string (nullable = true)

>>> ex6.show(5)
+-----+-----+-----+-----+-----+-----+
|placeid|name|food1|food2|food3|food4|placename|
+-----+-----+-----+-----+-----+-----+
|3|Joy|44|21|9|21|Food Town|
|2|Sam|13|17|6|8|Atlantic|
|5|Joe|21|7|12|8|Soup Bowl|
|3|Joe|37|19|10|36|Food Town|
|3|Jill|20|11|9|9|Food Town|
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

>>> |
```